Validation record, 2026-10-10: gated, human-approved multi-agent Claude workflow (sborycki-arch/voluntary-control-research)

Purpose

Validation of the stage-gated, human-approved, multi-agent Claude workflow as the subject of a methods paper, requested by Sean Borycki (program lead, sole approver). The research content it runs on (voluntary control of involuntary bodily functions, the 26-ability dataset) is a test-bed only, judged as stress-test material.

Basis

Anonymous clone of the public repository on 2026-10-10: main a690ad8ed40bf9e28299f33a1bf0fdddadf78dca, stage1-package 5d76f8a77a630220730a9f30caf8528537ef2e7b, review cf3b950ec3556e560fe8dbac47419b600638e2ba, transcripts bcf3b02df88cc7ba665d152f92835246252e1dfb; claude.ai export 5,084,660 bytes, sha256 f643aa45a6da51f317e5b50aff697489f44c3388ac2850d0073c9a745ccf30aa (verified). Not verifiable: Zenodo record 23199360, vcr-stage1-build.zip, the live tracker, and R (not installable by the documented route).

What the validation shows

Four validators ran on separate basis copies: A, B and C in round 1 on a broad task list; D in round 2, blind to A–C and the reconciliation, under the reproducibility schema the program lead set after reading round 1.

Validator	Model, effort	Tool calls	Duration	Verdict label
A	haiku, max	148	26.9 min	Does not validate yet
B	sonnet, medium	62	9.7 min	Validates with corrections
C	opus, low	60	13.0 min	Does not validate yet
D	sonnet, medium (clarified schema)	53	5.9 min	Does not validate yet

On the facts, all four agree (verified in each run): the deterministic tooling reproduced (trace logger 45/45, reconcile tool fail-closed, synthetic generator byte-identical, both manifests); the R pipeline was not reproduced (the documented apt route failed for every validator; no renv.lock; the committed 51-check log was read, not rerun); the Bayesian path never ran (bayesmeta absent in every log); no agent run was repeated (platform_default settings, single runs); no real-data result exists; and G1b, the program's own reproducibility test, has not run, is not blind (the 0.72 target sits in agent-loaded abilities.md:20) and cannot fail (the "or every difference traced" clause). None reports a reproduced R or real-data result.

On the verdict label, three of four said "Does not validate yet" (A, C, D). B said "Validates with corrections" on the broader round-1 task list.

Haiku, the lowest tier at maximum effort, was the most thorough: 148 tool calls, and it alone found the build's deletion of its own logger smoke-test record and the unfiled 13:37:42Z reviewer message. D reached the same facts in 5.9 min. Every validator made at least one error that the coordinator's reconciliation corrected.

Significance (coordinator judgment): the audit role gives consistent facts across three model tiers and effort levels; cross-model validator agreement is a candidate pre-specified outcome for the methods paper. Limits: A–C shared one briefing and the coordinator-prepared basis, D shared the basis, and one coordinator reconciled all four, so agreement partly reflects shared inputs. It is not evidence that the research results reproduce: the validators read the same record, and no real-data result exists to reproduce.

Verdict

Does not validate yet — reproducibility not demonstrated. The concept is not refuted; every component that could execute did. "Not refuted" is not "validated". Gates: two approval-channel deviations, Minor, record-keeping only; Sean's chat approvals (03:31:46Z, 11:45:02Z) preceded all work, the tracker followed 10 h 02 m 38 s and 1 h 49 m 22 s later, and no work preceded approval in any channel (verified).

Next steps (VALIDATION_PLAN_2026-10-10; nothing executed; thresholds set by Sean before each run)

- Phase 0: correct the method text (approval in any channel counts, quoted with timestamp and message ID; the tracker is the record), the trace schema and the record's execution errors.
- Phase 1: pinned, rebuildable R environment with bayesmeta; 51 structural plus value-level checks; built and run in two clean containers.
- Phase 2: at least three independent repeat agent runs per role; κ and exact-match thresholds; no adjudication escape.
- Phase 3: G1b run blind twice; target sealed by sha256 and removed from agent-loaded files; a pass rule that can fail.
- Phase 4: dry run executed twice; outputs committed with hashes; a count threshold replaces adjudication.
- Phase 5: cross-model replication of Phase 2 (and 3) on a second model tier.
- Phase 6: freeze; deposit with hash check; outcomes pre-specified in the frozen release; re-validation by two fresh validators of different tiers.

The program lead intends to pursue this full rerun for formal validation as time permits.

Limitations

The schema and verdict basis were set by the program lead after round 1; round-1 positions stay on record. Nobody checked R, the zip, Zenodo or the live tracker; thinking blocks are empty in the export. Validator independence is of analysis, not inputs.